

# Naïve Introduction to Word Embeddings

**Cristina España-Bonet**

UdS & DFKI, Saarbrücken, Germany

Summer Semester Seminar

17th April 2019

## *A word embedding is...*

**... a numerical representation of a word**

- Allow arithmetic operations on text

*Ex: time + flies*

## *A word embedding is...*

**... a numerical representation of a word**

- Allow arithmetic operations on text

*Ex: time + flies*

- Several names
  - Semantic Representations of Words
  - Word Vector Representations
  - Word Embeddings

## A word embedding is...

*time* = (1.844012, 0.590383, 1.003636, -0.577031, 1.515419, 1.097797, 1.812856, 0.933615, -2.396581, -0.931116, -0.719396, -0.376134, -1.204231, 0.045771, -0.287482, 1.084627, 4.399265, 1.516829, -0.838133, -1.881685, 0.108117, 2.345857, -1.292667, -2.286168, 3.419926, 4.260052, -1.016988, 3.140229, -3.161504, -0.800707, -1.433775, 2.290546, 1.932333, 0.714649, -3.033084, -0.958289, -1.704687, -1.597345, 1.525060, 3.337017, -2.787743, 1.479353, 3.452092, -3.242210, 0.532302, -0.551804, 2.344314, -0.919049, -1.872516, 0.080137, 1.208913, -2.136555, -2.218254, 0.206410, 0.133225, -1.521032, 1.735609, 2.885288, -2.048691, 2.375038, 0.316599, -0.254595, 2.159168, 1.118603, -0.775468, 0.933521, -0.351797, 2.193516, 2.499064, 2.818742, -0.213898, 0.446962, 1.767461, 1.342941, 1.117215, -0.042004, 4.199081, 3.041796, -1.770649, -0.528354, -2.067354, 0.283046, -0.099049, -0.105402, 2.823484, -2.583724, -2.906962, 0.592174, -3.029664, -0.170582, 0.406366, 1.963008, -3.229250, -3.499467, -0.136623, -1.551140, 0.348241, -1.597526, 0.703598, 3.122618, 0.466473, -0.113320, -2.119155, 1.092863, -0.908410, 0.253259, -1.082862, 4.408773, 2.419691, 2.343239, 0.703793, 1.270707, 0.410221, -1.293057, -0.799147, 2.214563, -0.212623, 1.206766, -0.731273, 2.308388, -1.029362, -2.080709, 0.749148, -1.412619, 1.073051, -2.498955, -0.520858, 1.391912, -1.181121, 1.523457, -1.245448, -0.290742, -2.589719, -0.366162, 3.586508, 0.908829, -1.125176, -0.937035, -1.163619, 1.759209, 3.678231, 0.019263, -0.395732, 1.142848, -0.500150, -3.005232, 2.287069, -0.524648, -0.944902, 0.038368, -1.093538, -0.697787, 0.767664, 2.399855, 2.425945, 1.563581, -1.086811, 0.372100, 1.400303, -2.278863, 0.643208, -0.459837, 1.756295, 2.057359, 3.140241, -1.740582, 1.386243, -1.822378, 1.528883, -1.984250, 1.214508, -1.336822, -0.321478, -0.162113, 0.272326, -2.673072, 0.612675, -0.657483, -0.557969, -3.358420, -2.559981, -1.683046, -1.314229, -2.425110, -2.506184, -1.606668, 1.332781, -2.760878, -2.400824, -1.806618, -2.406664, -1.169146, -1.838281, 0.588559, 2.285466, -0.401462, 1.632473, -0.510084, -2.072332, -2.627897, 2.531830, -2.524195, 2.035469, 1.906113, -1.257332, -4.039220, -0.467614, -2.275054, -3.409202, -0.014383, 0.445576, 1.461529, -1.318478, 0.061049, 0.280523, 2.173227, -0.027133, 2.791830, -0.728346, -1.804815, 1.245291, 0.970318, 2.646388, 0.246842, -1.823608, 1.888760, 0.265116, -2.027269, -0.089802, 0.389976, -0.654499, 2.565478, -2.647825, 2.658914, 1.385568, 2.306623, 0.476923, -0.869644, -0.170338, 0.495097, -2.604649, 0.610231, 0.739677, 0.322778, -2.042915, -1.353154, 0.177016, 1.840185, -0.271689, -0.401560, -0.421108, -0.185526, 1.041765, -4.599578, -0.829409, 0.076258, -0.503421, 1.891007, -0.931777, 0.434825, -0.467926, -1.417658, -0.320597, -4.084039, -3.899607, 0.977403, 0.774670, 3.269479, -1.031264, -0.433907, -2.305760, 0.811788, 2.347483, -1.254061, -0.861366, 0.080974, -3.666142, -0.363376, -2.384475, -4.290071, -0.924723, 1.257435, 1.223927, 0.276726, 1.541471, 1.274240, 1.883040, -1.987514, -0.809325, 1.252716, 1.812783, -0.511801, -1.657522, 1.196169, 0.804855, -1.861488, -2.113367, 0.429888, -0.920844, 0.377247)

## A word embedding is...

*flies* = (0.101159, 0.550446, 0.543801, -0.973852, -0.680835, 0.417193, -0.247181, 0.209725, -1.136055, -0.059531, -0.401640, 0.171540, 0.925121, -0.143815, 0.781714, -1.482425, 0.347008, -0.112342, 0.442418, -1.020457, -0.071752, 1.873548, -0.222886, -0.729569, -0.830224, -0.868407, 0.203496, 0.469911, -0.191363, 0.565102, 0.687738, 0.480823, 0.842358, -0.173656, -0.265585, 0.685740, 0.488047, -0.359772, -0.576064, -0.802884, 0.081554, 0.046882, -0.861532, -0.461855, 0.613098, -1.534642, -0.884534, 0.207728, 1.396512, -0.242900, -0.383959, 0.570844, -0.703350, -1.368813, -1.008194, 1.534660, 0.171693, 0.640925, -0.233116, 0.324685, 0.483171, 0.337947, -0.963290, -0.400558, 0.830977, 0.913474, 0.251693, -0.589420, -0.299622, 1.047515, -0.266679, -1.247186, 1.087610, -0.549028, 1.600710, -1.538516, -1.703301, -1.393499, -0.894448, 0.717204, 0.105767, -0.189234, -0.615609, -0.658315, 0.051877, 0.014180, -0.791282, 0.150424, 1.343751, -0.464859, 0.871426, 1.542864, -1.202150, -0.767113, -1.734738, 0.073633, -1.012583, 0.747787, 0.476070, -0.454807, 0.642685, -0.854152, -0.071798, 0.233724, 0.712329, -0.097752, -0.531132, 0.323271, -0.447342, 0.657913, 1.199492, -0.107360, -0.154234, -1.131168, 1.354793, 1.721385, -0.240023, 0.655765, -0.217006, -0.801722, 0.553369, 0.213377, 0.323267, -1.516051, 2.106244, -0.134282, 0.742155, 0.426344, 0.197991, -0.806768, 0.372546, -0.160200, -1.552847, -0.286178, -0.707796, 0.527352, -0.259658, 0.230387, 0.105294, -0.194481, 0.301772, -1.022163, 0.557191, 1.096709, 0.058422, -1.036384, 0.353412, -0.623097, -0.689515, 0.091472, 0.783885, 0.184088, -0.367950, 0.952462, 0.183704, 0.677562, 0.293917, -0.214309, -0.487794, 0.934296, 0.311513, 0.286514, -0.085511, 0.777691, 1.232603, -0.309367, -0.225086, 0.005091, -0.099195, -0.293117, 1.305563, 0.595816, 0.950316, 0.568706, -0.561446, 0.911634, -0.383941, 0.758054, -0.197820, 0.506777, -0.290767, -0.356727, 1.229474, -0.156489, -0.782741, -0.210163, -0.029169, 0.602664, 0.418375, 0.148975, -0.761796, 1.322690, -0.173410, 0.204111, -1.344531, 1.081905, -0.660543, -0.225615, -0.444753, -0.929671, 0.054136, 0.052031, -0.164926, 0.159312, -1.316333, 0.837011, -1.290353, 0.958403, 1.247478, 0.442009, 0.455497, -1.856268, -0.358823, -0.230839, -0.206271, 0.227012, -0.454163, 0.747798, -1.252855, 1.436849, -0.427915, -0.810428, -0.628144, -0.288458, 0.087355, 0.356739, 0.153036, 0.516594, -0.504978, 0.814432, 1.052940, 1.094526, -0.219595, 0.722178, 0.267325, -0.087458, -1.270262, -0.039461, 0.991926, -0.112005, -0.009605, 0.149920, 0.164717, 0.280475, 0.966384, 0.327598, 0.189590, -0.208946, 0.838261, 0.051847, -0.277932, -0.788527, -0.768702, -1.688721, 0.388215, 0.170153, -0.555723, -0.529565, -0.528982, -0.659930, 0.588041, -0.368195, -0.850188, -0.004996, 0.925476, 1.046587, -0.731761, 0.519435, 0.193188, -0.709557, 0.123329, -0.454316, 1.885830, -0.201841, -0.728933, -0.953455, -0.205837, -0.724068, 0.120158, 1.765389, -0.192159, 1.062490, -0.002634, 0.125790, -0.846565, 0.548899, -1.062821, -2.146826, 0.134681, 0.570950, 0.851783, 0.436544, 0.688986, 1.229008, 1.435449, 0.118766, -0.132411, 2.527890, 0.778142, 0.269093)

## *The Purpose of the First Day is...*

**...to answer several questions:**

- How can we obtain those numbers?
- What's word2vec?
- Is it the only way to obtain those numbers?
- Do the vectors (and components!) have any semantic meaning?
- Are we crazy by summing or multiplying words?

# Outline

- 1 Introduction
- 2 Frequency-based Embeddings
- 3 Prediction-based Embeddings
- 4 Beyond *Word* Embeddings
- 5 Software & References

# Introduction

- 1 Introduction
  - Distributional Hypothesis
  - Term Frequencies
- 2 Frequency-based Embeddings
- 3 Prediction-based Embeddings
- 4 Beyond *Word* Embeddings
- 5 Software & References



# Introduction

## *Distributional Hypothesis, Contextuality*

**Never ask for the meaning of a word in isolation, but  
only in the context of a sentence**  
(Frege, 1884)

# Introduction

## *Distributional Hypothesis, Contextuality*

**Never ask for the meaning of a word in isolation, but only in the context of a sentence**  
(Frege, 1884)

**For a large class of cases... the meaning of a word is its use in the language**  
(Wittgenstein, 1953)

**You shall know a word by the company it keeps**  
(Firth, 1957)

**Words that occur in similar contexts  
tend to have similar meaning**

**(Harris, 1954)**

# Introduction

## *Similar Meanings...*

- ...need for a concept of **distance** to be defined.
- **Geometry** is the branch of mathematics that deals with distances
- **Vector spaces** and linear algebra are our tools

# Introduction

*Similar Meanings...*

Sumo

Basketball


Football

Cherry

Strawberry

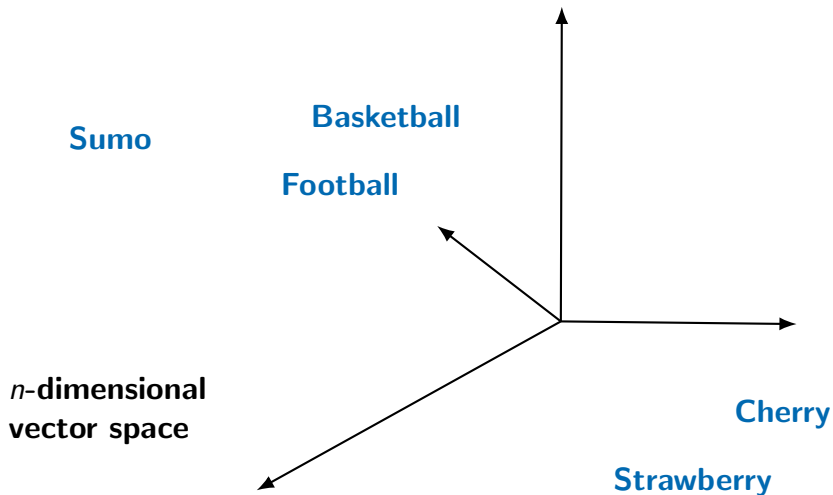
# Introduction

## *Similar Meanings...*



# Introduction

## *Word Vector Space*



# Introduction

## *How to Obtain a Vector for a Word?*

Naïve example: **term frequencies** in a corpus

- The basis in our vector space is the vocabulary of the corpus
- Consider the document in which a word occurs its context
- Each word is characterised as the number of times it appears in each document



# Introduction

## *Example: Toy Corpus*

S1: We like to play some sport in the afternoon, I like basketball but John likes sumo more.

S2: Sumo is a sport where a rikishi attempts to force another wrestler out of a circular ring.

S3: Messi scored 4 goals yesterday and kept the ball as a memory of this fantastic sports afternoon!

S4: I ate too many cherries yesterday.

**Vocabulary:** {*like, play, sport, afternoon, basketball, John, ball, sumo, rikishi, attempt, force, werstler, circular, ring, Messi, score, goal, yesterday, keep, memory, fantastic, eat, cherry*}

# Introduction

## Example: Occurrence Matrix

	like	play	sport	afternoon	basketball	John	ball	sumo	rikishi	attempt	force	werstler	circular	ring	Messi	score	goal	yesterday	keep	memory	fantastic	eat	cherry
<i>S1</i>	2	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
<i>S2</i>	0	0	1	0	0	0	0	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0
<i>S3</i>	0	0	1	1	0	0	1	0	0	0	0	0	0	0	1	1	1	1	1	1	1	0	0
<i>S4</i>	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	1	1

# Introduction


## Example: Occurrence Matrix

	like	play	sport	afternoon	basketball	John	ball	sumo	rikishi	attempt	force	werstler	circular	ring	Messi	score	goal	yesterday	keep	memory	fantastic	eat	cherry
S1	2	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
S2	0	0	1	0	0	0	0	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0
S3	0	0	1	1	0	0	1	0	0	0	0	0	0	0	1	1	1	1	1	1	1	0	0
S4	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	1	1

document vector

# Introduction

## Example: Occurrence Matrix



	like	play	sport	afternoon	basketball	John	ball	sumo	rikishi	attempt	force	werstler	circular	ring	Messi	score	goal	yesterday	keep	memory	fantastic	eat	cherry
S1	2	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
S2	0	0	1	0	0	0	0	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0
S3	0	0	1	1	0	0	1	0	0	0	0	0	0	0	1	1	1	1	1	1	1	0	0
S4	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	1	1

word vector

# Introduction

## *Example: Text Similarity*

**Euclidean distance**       $d(\vec{x}, \vec{y}) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$

basketball  $\rightarrow \{1, 0, 0, 0\}$

sumo  $\rightarrow \{1, 1, 0, 0\}$

cherry  $\rightarrow \{0, 0, 0, 1\}$

$$d(\text{basketball}, \text{sumo}) = \sqrt{(1-1)^2 + (0-1)^2 + (0-0)^2 + (0-0)^2} = 1$$

$$d(\text{basketball}, \text{cherry}) = \sqrt{(1-0)^2 + (0-1)^2 + (0-0)^2 + (0-0)^2} = \sqrt{2}$$

$$d(\text{sumo}, \text{cherry}) = \sqrt{(1-0)^2 + (1-0)^2 + (0-0)^2 + (0-1)^2} = \sqrt{3}$$

$$d(\text{basketball}, \text{sumo}) < d(\text{basketball}, \text{cherry}) < d(\text{sumo}, \text{cherry})$$

# Introduction

## *Example: Text Similarity*

**Cosine similarity**       $\text{sim}(\vec{x}, \vec{y}) = \frac{\sum_{i=1}^n x_i y_i}{|\vec{x}| |\vec{y}|}$

basketball  $\rightarrow \{1, 0, 0, 0\}$

sumo  $\rightarrow \{1, 1, 0, 0\}$

cherry  $\rightarrow \{0, 0, 0, 1\}$

$\text{sim}(\text{basketball}, \text{sumo})=1$


$\text{sim}(\text{basketball}, \text{cherry})=0$

$\text{sim}(\text{sumo}, \text{cherry})=0$

$\text{sim}(\text{basketball}, \text{sumo}) > \text{sim}(\text{basketball}, \text{cherry}) = \text{sim}(\text{sumo}, \text{cherry})$


# Introduction

## *Similarity vs. Distance*



# Introduction

## *Similarity vs. Distance*





# Frequency-based Embeddings

- 1 Introduction
- 2 Frequency-based Embeddings
  - TF-IDF
  - Co-Occurrence
- 3 Prediction-based Embeddings
- 4 Beyond *Word* Embeddings
- 5 Software & References

# Frequency-based Embeddings

## *Frequency-based Embeddings*

- *Term frequency word vectors*
- TF-IDF word vectors
- Co-occurrence word vectors

# Frequency-based Embeddings

## *Term Frequency-Inverse Document Frequency, TF-IDF*

### **Term Frequency**

How frequently a term occurs in a document  $d$  normalised to account for  $d$  length

$$\text{TF}(t, d) = \frac{\text{Number of times term } t \text{ appears in a document } d}{\text{Total number of terms in } d}$$

# Frequency-based Embeddings

## *Term Frequency-Inverse Document Frequency, TF-IDF*

### **Term Frequency**

How frequently a term occurs in a document  $d$  normalised to account for  $d$  length

$$\text{TF}(t, d) = \frac{\text{Number of times term } t \text{ appears in a document } d}{\text{Total number of terms in } d}$$

### **Inverse Document Frequency**

Measures how important a term is (low weight for stop words)

$$\text{IDF}(t, D) = \log_e \left( \frac{\text{Total number of documents } D}{\text{Number of documents with term } t \text{ in it}} \right)$$

# Frequency-based Embeddings

*Term Frequency-Inverse Document Frequency, TF-IDF*

Trivially...

$$\text{TF-IDF}(t, d, D) = \text{TF}(t, d) \times \text{IDF}(t, D)$$

# Frequency-based Embeddings

## *Example: Toy Corpus*

*d1*: We like to play some sport in the afternoon, I like basketball but John likes sumo more.

*d2*: Sumo is a sport where a rikishi attempts to force another wrestler out of a circular ring.

*d3*: Messi scored 4 goals yesterday and kept the **ball** as a memory of this fantastic sports afternoon!

*d4*: I ate too many cherries yesterday.

$$\mathbf{TF}(\mathbf{ball}) = \left(0, 0, \frac{1}{17}, 0\right);$$

# Frequency-based Embeddings

## *Example: Toy Corpus*

*d1:* We like to play some sport in the afternoon, I like basketball but John likes sumo more.

*d2:* Sumo is a sport where a rikishi attempts to force another wrestler out of a circular ring.

*d3:* Messi scored 4 goals yesterday and kept the **ball** as a memory of this fantastic sports afternoon!

*d4:* I ate too many cherries yesterday.

$$\mathbf{TF}(\mathbf{ball}) = \left(0, 0, \frac{1}{17}, 0\right); \quad \mathbf{IDF}(\mathbf{ball}) = \log_e \left(\frac{4}{1}\right);$$

# Frequency-based Embeddings

## *Example: Toy Corpus*

*d1:* We like to play some sport in the afternoon, I like basketball but John likes sumo more.

*d2:* Sumo is a sport where a rikishi attempts to force another wrestler out of a circular ring.

*d3:* Messi scored 4 goals yesterday and kept the **ball** as a memory of this fantastic sports afternoon!

*d4:* I ate too many cherries yesterday.

$$\mathbf{TF}(\mathbf{ball}) = \left(0, 0, \frac{1}{17}, 0\right); \quad \mathbf{IDF}(\mathbf{ball}) = \log_e \left(\frac{4}{1}\right);$$

$$\mathbf{TF-IDF}(\mathbf{ball})_3 = \frac{1}{17} \times \log_e(4) = 0.08$$



# Frequency-based Embeddings

## *Example: Toy Corpus*

*d1*: We like to play some sport in the afternoon, I like basketball but John likes sumo more.

*d2*: Sumo is **a** sport where **a** rikishi attempts to force another wrestler out of **a** circular ring.

*d3*: Messi scored 4 goals yesterday and kept the ball as **a** memory of this fantastic sports afternoon!

*d4*: I ate too many cherries yesterday.

$$\mathbf{TF}(\mathbf{a}) = \left(0, \frac{3}{17}, \frac{1}{17}, 0\right);$$

# Frequency-based Embeddings

## *Example: Toy Corpus*

*d1:* We like to play some sport in the afternoon, I like basketball but John likes sumo more.

*d2:* Sumo is **a** sport where **a** rikishi attempts to force another wrestler out of **a** circular ring.

*d3:* Messi scored 4 goals yesterday and kept the ball as **a** memory of this fantastic sports afternoon!

*d4:* I ate too many cherries yesterday.

$$\mathbf{TF(a)} = \left(0, \frac{3}{17}, \frac{1}{17}, 0\right); \quad \mathbf{IDF(a)} = \log_e \left(\frac{4}{2}\right);$$

# Frequency-based Embeddings

## *Example: Toy Corpus*

*d1:* We like to play some sport in the afternoon, I like basketball but John likes sumo more.

*d2:* Sumo is **a** sport where **a** rikishi attempts to force another wrestler out of **a** circular ring.

*d3:* Messi scored 4 goals yesterday and kept the ball as **a** memory of this fantastic sports afternoon!

*d4:* I ate too many cherries yesterday.

$$\mathbf{TF(a)} = \left(0, \frac{3}{17}, \frac{1}{17}, 0\right); \quad \mathbf{IDF(a)} = \log_e \left(\frac{4}{2}\right);$$

$$\mathbf{TF-IDF(a)}_2 = \frac{3}{17} \times \log_e(2) = 0.12; \quad \mathbf{TF-IDF(a)}_3 = 0.04$$

# Frequency-based Embeddings

*Term Frequency-Inverse Document Frequency, TF-IDF*

- Word vectors of  $D$  dimensions
- Distances between words as before:
  - Euclidean distance
  - Cosine similarity
  - ...

# Frequency-based Embeddings

## *Co-Occurrence Matrix, Count Vectors*

- Words co-occurrence statistics describes how words occur together
- Counts how two or more words occur together in a given corpus

# Frequency-based Embeddings

## *Example: Toy Corpus*

*d1*: We like to play some sport in the afternoon, I like basketball but John likes sumo more.

*d2*: Sumo is a sport where a rikishi attempts to force another wrestler out of a circular ring.

*d3*: Messi scored 4 goals yesterday and kept the ball as a memory of this fantastic sports afternoon!

*d4*: I ate too many cherries yesterday.

# Frequency-based Embeddings

## *Example: Toy Corpus*

*d1*: We **like** to play some sport in the afternoon, I **like** basketball but John likes sumo more.

*d2*: Sumo is **a** sport where **a** rikishi **attempts** to force another wrestler out of **a** circular ring.

*d3*: **Messi** scored 4 goals yesterday and kept the **ball** as **a** memory of this fantastic sports afternoon!

*d4*: I ate too many cherries yesterday.

# Frequency-based Embeddings

### Example: Co-Occurrence Matrix

[illegible]



# Frequency-based Embeddings

*Co-occurrence matrix, count vectors*

- Simple bigram frequencies of all possible word-pairs need a size  $N \times N$  matrix to represent  $N$  words in a corpus
- Real models use context windows, not only bigrams
- Counts are converted into probabilities
- In general, one has sparse matrices
- Dimensionality reduction (SVD, see next session)

# Frequency-based Embeddings

## Aside Comment: One-Hot Encodings

	like	to	sport	afternoon	basketball	John	ball	a	rikishi	attempt	force	wrestler	circular	as	Messi	score	goal	yesterday	keep	memory
<i>like</i>	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
<i>attempt</i>	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0
<i>Messi</i>	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0
<i>a</i>	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0
...																				

One-hot encoding for *like* with this vocabulary:

*like* = (1, 0)

# Frequency-based Embeddings

*Aside Comment: One-Hot Encodings*

	like	to	sport	afternoon	basketball	John	ball	a	rikishi	attempt	force	wrestler	circular	as	Messi	score	goal	yesterday	keep	memory
<i>like</i>	0	1	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
<i>attempt</i>	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
<i>Messi</i>	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0
<i>a</i>	0	0	1	0	0	0	0	0	1	0	0	0	1	0	0	0	0	0	0	1
...																				

Co-occurrence word vector for *like* with this vocabulary in the previous corpus:

*like* = (0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0)

# Prediction-based Embeddings

- 1 Introduction
- 2 Frequency-based Embeddings
- 3 Prediction-based Embeddings**
  - Continuous Bag of Words
  - Skip-Gram Model
  - Demos
- 4 Beyond *Word* Embeddings
- 5 Software & References

# Prediction-based Embeddings

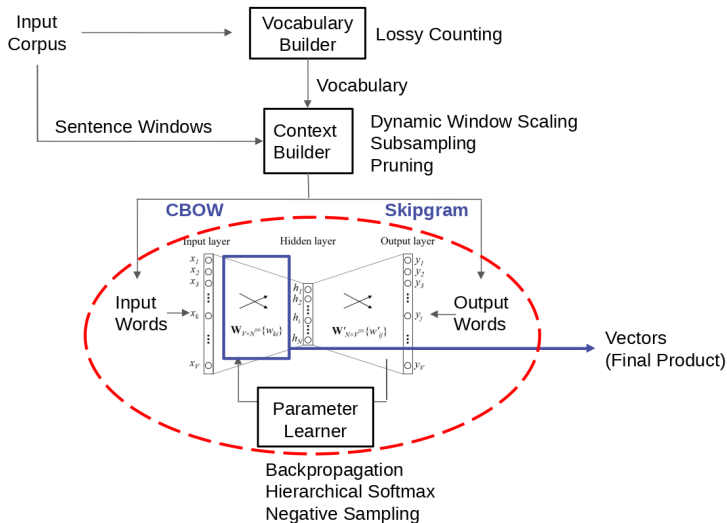
## *Word Embeddings (word2vec example)*

Word vectors learned by a neural network in two tasks:

- 1 predict the probability of a **word given a context**  
(CBow)
- 2 predict the **context given a word**  
(skip-gram)

# Prediction-based Embeddings


## Word Embeddings (word2vec example)



Credits: Xin Rong

# Prediction-based Embeddings

## *Word Embeddings (word2vec example)*



Look at the network: simple feed-forward network learned by backpropagation with cross-entropy loss

No deep learning at all!

# Prediction-based Embeddings

## *Word Embeddings (word2vec example)*

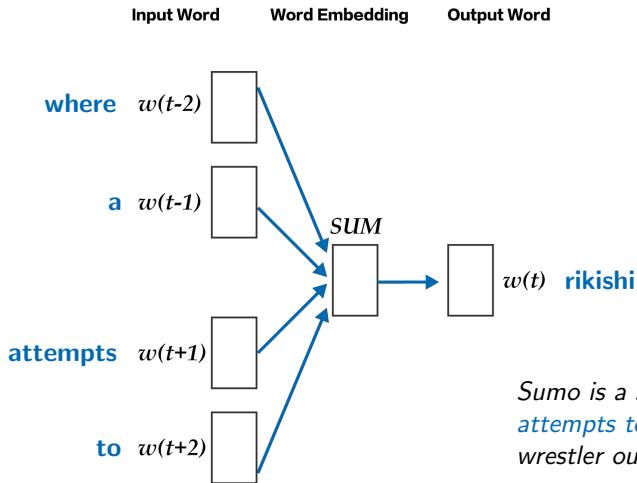
### **Comments:**

- A hidden layer in a NN interprets the input in his own way to optimise his work in the concrete task
- The size of the hidden layer gives you the dimension of the word embeddings
- Too few neurons could not have enough capacity to learn everything needed
- Too many neurons would need a very large corpus to be meaningful



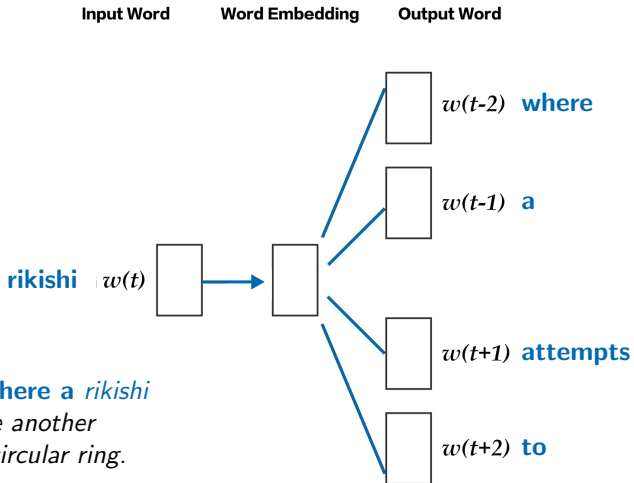
# Prediction-based Embeddings

## Continuous Bag of Words, CBoW



# Prediction-based Embeddings

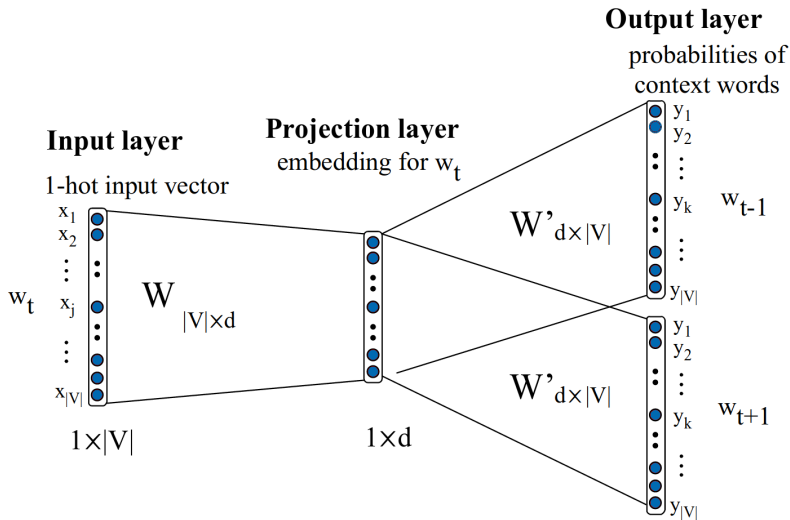
## Skip-Gram Model



*Sumo is a sport **where a rikishi attempts to** force another wrestler out of a circular ring.*

# Prediction-based Embeddings

## More Detailed Architecture (skip-gram)



# Prediction-based Embeddings

*More Detailed Architecture (schematic matrix visualisation)*

$$\begin{matrix} \begin{pmatrix} V \\ \vdots \\ \vdots \end{pmatrix} & \begin{pmatrix} V \times d \\ \vdots \\ \vdots \end{pmatrix} & \begin{pmatrix} d \\ \vdots \\ \vdots \end{pmatrix} & \begin{pmatrix} d \times V \\ \vdots \\ \vdots \end{pmatrix} & \begin{pmatrix} V \\ \vdots \\ \vdots \end{pmatrix} \\ \mathbf{x} & \mathbf{W} & \mathbf{h} & \mathbf{W}' & \mathbf{y} \end{matrix}$$

## Input Embedding

The row  $i$  of the input matrix  $\mathbf{W}$  is the  $1 \times d$  for word  $i$  in the vocabulary

# Prediction-based Embeddings

*More Detailed Architecture (schematic matrix visualisation)*

$$\begin{pmatrix} V \\ \end{pmatrix} \begin{pmatrix} V \times d \\ \end{pmatrix} \begin{pmatrix} d \\ \end{pmatrix} \begin{pmatrix} d \times V \\ \end{pmatrix} \begin{pmatrix} V \\ \end{pmatrix}$$

$\mathbf{x} \qquad \mathbf{W} \qquad \mathbf{h} \qquad \mathbf{W'} \qquad \mathbf{y}$

## Output Embedding

The column  $j$  of the output matrix  $\mathbf{W'}$  is the  $d \times 1$  for word  $j$  in the vocabulary

# Prediction-based Embeddings

*Observations (Tensorflow Tutorial)*

## CBoW

- Smoothes over a lot of the distributional information by treating an entire context as one observation. This turns out to be a useful thing for **smaller datasets**



## Skip-gram

- Treats each context-target pair as a new observation, and this tends to do better when we have **larger datasets**

# Prediction-based Embeddings

## *Nice Properties*


King - Man + Woman = Queen



(Mikolov et al., NAACL HLT, 2013)

# Prediction-based Embeddings


## Nice Properties





# Prediction-based Embeddings

## Nice Properties




*2D tSNE projection of the main characters of Game of thrones colored by House*

# Prediction-based Embeddings

## Nice Properties

(Luong, Pham & Manning, NAACL, 2015)



*Barnes-Hut-SNE visualisation of bilingual  
embeddings German/English*

# Beyond *Word* Embeddings

- 1 Introduction
- 2 Frequency-based Embeddings
- 3 Prediction-based Embeddings
- 4 Beyond *Word* Embeddings
  - Compositionality
  - Sentence and Document Embeddings
- 5 Software & References

# Beyond *Word* Embeddings

*Is Language Compositional?*

**The meaning of a compound expression is a function of the meanings of its parts and of the way they are syntactically combined.**

**(Partee, 1984)**

# Beyond *Word* Embeddings

*Is Language Compositional?*

Computer Scientist-like Background Yes!

$$\begin{aligned} &\text{meaning}(\text{Eat the icecream}) = \\ &\text{meaning}(\text{Eat}) + \text{meaning}(\text{the}) + \text{meaning}(\text{icecream}) \end{aligned}$$

# Beyond *Word* Embeddings

*Is Language Compositional?*

Computer Scientist-like Background Yes!

$$\begin{aligned} &\text{meaning}(\text{Eat the icecream}) = \\ &\text{meaning}(\text{Eat}) + \text{meaning}(\text{the}) + \text{meaning}(\text{icecream}) \end{aligned}$$

Linguist-like Background No!

$$\begin{aligned} &\text{meaning}(\text{Break the ice}) \neq \\ &\text{meaning}(\text{Break}) + \text{meaning}(\text{the}) + \text{meaning}(\text{ice}) \end{aligned}$$

# Beyond *Word* Embeddings

*Is Language Compositional?*

**Computer Scientist-like Background** Yes!

$$\begin{aligned} &\text{meaning}(\text{Eat the icecream}) = \\ &\text{meaning}(\text{Eat}) + \text{meaning}(\text{the}) + \text{meaning}(\text{icecream}) \end{aligned}$$

**Linguist-like Background** No!

$$\begin{aligned} &\text{meaning}(\text{Break the ice}) \neq \\ &\text{meaning}(\text{Break}) + \text{meaning}(\text{the}) + \text{meaning}(\text{ice}) \end{aligned}$$

# Beyond *Word* Embeddings

## *Representations for Phrases, Sentences or Paragraphs*

- **Composition** of word embeddings using operations  $(+, \times)$  on vectors and matrices
- **Latent paragraph vectors** in word2vec-like NNs
- Internal representations in **seq2seq** architectures or auto-encoders (NMT context vectors, skip-thought vectors...)
- **Contextual** word vectors (LM by-products for instance)



# Beyond *Word* Embeddings

## Composition I

(Mitchell & Lapata, 2010)

Table 5

Composition functions considered in our experiments

Model	Function
Additive	$p_i = u_i + v_i$
Kintsch	$p_i = u_i + v_i + n_i$
Multiplicative	$p_i = u_i \cdot v_i$
Tensor product	$p_{i,j} = u_i \cdot v_j$
Circular convolution	$p_i = \sum_j u_j \cdot v_{i-j}$
Weighted additive	$p_i = \alpha v_i + \beta u_i$
Dilation	$p_i = v_i \sum_j u_j u_j + (\lambda - 1) u_i \sum_j u_j v_j$
Head only	$p_i = v_i$
Target unit	$p_i = v_i(t_1 t_2)$

# Beyond *Word* Embeddings

## Composition II

(Mitchell & Lapata, 2010)

Table 6

Correlation coefficients of model predictions with subject similarity ratings (Spearman's  $\rho$ ) using a simple semantic space

Model	Adjective–Noun	Noun–Noun	Verb–Object
Additive	.36	.39	.30
Kintsch	.32	.22	.29
Multiplicative	.46	.49	.37
Tensor product	.41	.36	.33
Convolution	.09	.05	.10
Weighted additive	.44	.41	.34
Dilation	.44	.41	.38
Target unit	.43	.34	.29
Head only	.43	.17	.24
Humans	.52	.49	.55


# Beyond *Word* Embeddings

## *Distributed Representations of Sentences and Documents*

- word2vec-like architecture where a **document vector** is added to word vectors and learned simultaneously (PV-DM)
- **Generation** of words from a document from its document vector (skip-gram-like architecture!) (PV-DBoW)

# Beyond *Word* Embeddings

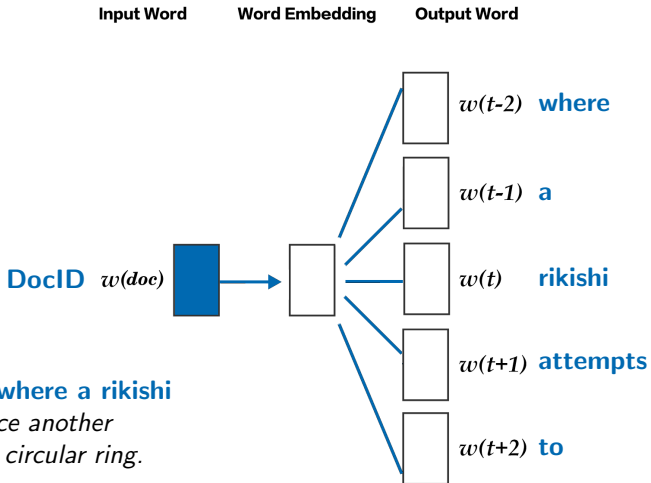
## Distributed Memory Model of Paragraph Vectors, PV-DM



Sumo is a sport *where a rikishi attempts* to force another wrestler out of a circular ring.


# Beyond Word Embeddings

## Distributed Bag of Words version of Paragraph Vector



# Beyond *Word* Embeddings

*Language Model as Complementary Task for Contextual Vectors*



<http://jalamar.github.io/illustrated-bert/>

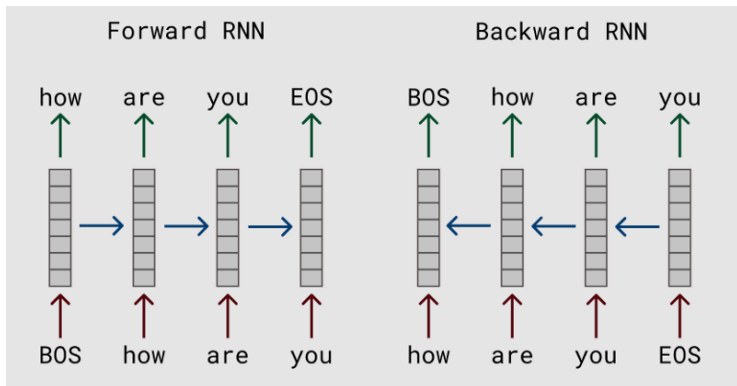
# Beyond *Word* Embeddings

## *Language Model as Complementary Task for Contextual Vectors*

- Deep **contextualised** word representations
- Mainly based on LSTMs or transformer architectures
- Example: ELMo, Embeddings from Language Models
- **Bidirectional language models** with LSTMs

# Beyond *Word* Embeddings

*Language Model as Complementary Task for Contextual Vectors*




[https://medium.com/@plusepsilon/  
the-bidirectional-language-model-1f3961d1fb27](https://medium.com/@plusepsilon/the-bidirectional-language-model-1f3961d1fb27)



# Beyond *Word* Embeddings

*Language Model as Complementary Task for Contextual Vectors*



[https://medium.com/@plusepsilon/  
the-bidirectional-language-model-1f3961d1fb27](https://medium.com/@plusepsilon/the-bidirectional-language-model-1f3961d1fb27)

# Beyond *Word* Embeddings

## *Language Model as Complementary Task for Contextual Vectors*

Word vectors and contextual word vectors are used differently:

- **Word vectors**: dictionary look-up of words and their corresponding vectors, they are static entities
  - Good to initialise input word embeddings in several NLP tasks

# Beyond *Word* Embeddings

## *Language Model as Complementary Task for Contextual Vectors*


Word vectors and contextual word vectors are used differently:

- **Word vectors**: dictionary look-up of words and their corresponding vectors, they are static entities
  - Good to initialise input word embeddings in several NLP tasks
- **Contextual word vectors**: vectors on-the-fly by passing text through a deep learning model, they would only be static if we could generate all sentences in a language!
  - Good for transfer learning into several NLP tasks (SotA in lots of tasks!)


# Beyond Word Embeddings

## Seq2Seq Internal Representations: (Multilingual) NMT

decoder



encoder




dense vectors

one-hot encoding

# Beyond *Word* Embeddings

## Multilingual Semantic Space for NMT Context Vectors

(Esaña-Bonet & van Genabith, 2018)



ML-NMT  $\{de, en, nl, it, ro\} \rightarrow \{de, en, nl, it, ro\}$  with TED talks

# Summary

- 1 Introduction
- 2 Frequency-based Embeddings
- 3 Prediction-based Embeddings
- 4 Beyond *Word* Embeddings
- 5 Software & References

# Summary

## *Assuming Contextuality Holds...*

- A word can be represented by a vector that describes it with respect to the context it is usually used
- This vector can be estimated by **counts** in a corpus
- This vector can be **learned** by examples in a corpus
- In fact, skipgram and co-occurrence matrix factorisation are equivalent under certain conditions
- They are very useful to characterise text, but we haven't talked about **ambiguity** problems for instance

# Summary

## *Assuming Compositionality Holds...*

- A sentence can be decomposed into the vectors of its constituents
- Simple operations such as **sum** and **product** work surprisingly well
- Sentence/Document vectors can also be **learned** in more general tasks such as translation
- They are very useful to characterise text, but we haven't talked about **idioms** or **negation** problems for instance



# Software & References

- 1 Introduction
- 2 Frequency-based Embeddings
- 3 Prediction-based Embeddings
- 4 Beyond *Word* Embeddings
- 5 Software & References

# Software & References

## *Libraries & Packages*

- **word2vec**

`https://github.com/dav/word2vec`

- **fastText**

`https://github.com/facebookresearch/fastText`

- **Gensim**

`https://radimrehurek.com/gensim/models/word2vec.html`

- **GloVe**

`https://nlp.stanford.edu/projects/glove/`

# Software & References

## *Libraries & Packages*

- **ELMo**

`https://github.com/allenai/allennlp/blob/master/allennlp/modules/elmo.py`

- **Open AI**

`https://github.com/openai/gpt-2`

- **BERT**

`https://github.com/google-research/bert`

- **InferSent**

`https://github.com/facebookresearch/InferSent`

# Software & References

## *Basic References*

- Stephen Clark. 2015. **Vector Space Models of Lexical Meaning**. *Handbook of Contemporary Semantic Theory*. Second edition, edited by Shalom Lappin and Chris Fox. Chapter 16. Pages 493–522. Wiley-Blackwell.
- Jeff Mitchell and Mirella Lapata. 2010. **Composition in distributional models of semantics**. *Cognitive Science*, 34(8). Pages 1388–1429.
- Daniel Jurafsky and James H. Martin. 2017. **Speech and Language Processing**. Chapter 16: Semantics with Dense Vectors.
- Xin Rong. 2015. **word2vec Parameter Learning Explained**. *eprint arXiv:1411.2738*.

# Software & References

## Basic References II

- Tomas Mikolov, Kai Chen, Greg Corrado and Jeffrey Dean. 2013. **Efficient Estimation of Word Representations in Vector Space**. *Proceedings of the Workshop at International Conference on Learning Representations (ICLR)*. Pages 1–12.
- Jeffrey Pennington, Richard Socher, and Christopher D. Manning. 2014. **GloVe: Global Vectors for Word Representation**. *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Pages 1532–1543
- Quoc V. Le and Tomas Mikolov. 2014. **Distributed Representations of Sentences and Documents**. *Proceedings of the 31st International Conference on Machine Learning (ICML)*, in PMLR 32(2). Pages 1188–1196

# Software & References

## Basic References III

- Matthew E. Peters, Mark Neumann, Mohit Iyyer, Matt Gardner, Christopher Clark, Kenton Lee, Luke Zettlemoyer. 2018. **Deep contextualized word representations**. *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*. Pages 2227–2237.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, Kristina Toutanova. 2018. **BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding**. <https://arxiv.org/abs/1810.04805>.
- Alan Akbik, Duncan Blythe, Roland Vollgraf. 2018. **Contextual String Embeddings for Sequence Labeling**. *Proceedings of the 27th International Conference on Computational Linguistic*). Pages 1638–1649.

Everything clear? Too much information?

**Questions?**

# Naïve Introduction to Word Embeddings

**Cristina España-Bonet**

UdS & DFKI, Saarbrücken, Germany

Summer Semester Seminar

17th April 2019



# Extra Slides

## *Singular-Value Decomposition, SVD*

- Linear algebra

# Extra Slides

## *Singular-Value Decomposition, SVD*

- Linear algebra
- **Factorisation** of a matrix  $\mathbf{M}$  as  $\mathbf{M} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$

- Linear algebra
- **Factorisation** of a matrix  $\mathbf{M}$  as  $\mathbf{M} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$ 
  - ✓  $\mathbf{U}$  is an  $m \times m$  orthogonal matrix,


- Linear algebra
- **Factorisation** of a matrix **M** as  **$M = U\Sigma V^T$** 
  - ✓ **U** is an  $m \times m$  **orthogonal matrix**,
    - **$U^T U = U U^T = I$**
    - or, equivalently,  **$U^T = U^{-1}$**

- Linear algebra
- **Factorisation** of a matrix  $\mathbf{M}$  as  $\mathbf{M} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$ 
  - ✓  $\mathbf{U}$  is an  $m \times m$  orthogonal matrix,
  - ✓  $\mathbf{\Sigma}$  is a diagonal  $m \times n$  matrix with non-negative real numbers,

- Linear algebra
- **Factorisation** of a matrix  $\mathbf{M}$  as  $\mathbf{M} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$ 
  - ✓  $\mathbf{U}$  is an  $m \times m$  orthogonal matrix,
  - ✓  $\mathbf{\Sigma}$  is a diagonal  $m \times n$  matrix with non-negative real numbers,
  - ✓  $\mathbf{V}^T$  is the conjugate transpose of an  $n \times n$  orthogonal matrix

# Extra Slides

## SVD: $2 \times 2$ Geometric Interpretation



a linear transformation is a rotation or reflection, followed by a scaling, followed by another rotation or reflection

# Extra Slides

## *Singular-Value Decomposition, SVD*

$$\begin{pmatrix} m \times n \end{pmatrix} = \begin{pmatrix} m \times m \end{pmatrix} \begin{pmatrix} m \times n \end{pmatrix} \begin{pmatrix} n \times n \end{pmatrix}$$

**M**

**=**

**U**

**Σ**

**V<sup>T</sup>**



# Extra Slides

## *SVD: Singular Values*

$$\mathbf{\Sigma} = \begin{pmatrix} \sigma_1 & & & \\ & \ddots & & \\ & & 0 & \\ & & & \ddots \\ 0 & & & & \sigma_r \\ & & & & & 0 \end{pmatrix};$$

$\sigma_1 \dots \sigma_r$ , singular values of  $\mathbf{M}$  (in decreasing order)

$r$ , rank of  $\mathbf{M}$

# Extra Slides

## SVD: Singular Values


$$\mathbf{\Sigma} = \begin{pmatrix} \sigma_1 & & & \\ & \ddots & & \\ & & 0 & \\ & & & \ddots \\ & 0 & & & \sigma_r \\ & & & & & 0 \end{pmatrix}; \quad \mathbf{M}_r = \sum_{i=1}^r \sigma_i \vec{u}_i \vec{v}_i^T$$

$\sigma_1 \dots \sigma_r$ , singular values of  $\mathbf{M}$  (in decreasing order)

$r$ , rank of  $\mathbf{M}$

# Extra Slides

## *SVD: Application, Latent Semantic Analysis*



$$\Sigma_r \Rightarrow M_r$$

<https://nlp.stanford.edu/IR-book/pdf/18lsi.pdf>

Christopher D. Manning, Prabhakar Raghavan, and Hinrich Schütze. 2008. **Introduction to Information Retrieval**. Cambridge University Press, New York, NY, USA.